

GTSC2143 Machine Learning for Business

Tutorial 4

Please write down your answers in this document and submit it at iSpace by the end of this tutorial.

Activity 1. Data Loading and Preprocessing

1. Load the Dataset

- a) Load the house prices dataset using pandas:

```
1. import pandas as pd
2. import numpy as np
3. from sklearn.model_selection import train_test_split
4. from sklearn.linear_model import LinearRegression
5. from sklearn.metrics import mean_squared_error, r2_score
6. import matplotlib.pyplot as plt
7. import seaborn as sns
8.
9. # Load the house prices dataset
10. data = pd.read_csv("GTSC2143-Lecture 4 predicting-house-prices-assignment_home_data.csv")
```

- b) Display basic information:

- Dataset shape
- First 5 rows
- Column names and data types

- c) Check for any missing values in the dataset

2. Data Filtering

- a) Filter the data to exclude the house whose id is '1925069082'

- b) Display the shape of the filtered dataset

3. Train/Test Split

- a) Split the filtered data into training (80%) and testing (20%) sets using `train_test_split`
- b) Use `random_state=42` for reproducible results
- c) Display the shapes of training and testing sets
- d) Analysis: Write 2-3 sentences explaining the purpose of train/test split in machine learning.

Activity 2. Model Training, Evaluation and Prediction

1. Feature Selection and Model Training

- a) Select these features for your regression model:

- 'bedrooms'
- 'bathrooms'
- 'sqft_living'
- 'sqft_lot'
- 'floors'
- 'zipcode'

- b) Create and train a Linear Regression model using the selected features

- c) Display the model coefficients and intercept

- d) Analysis: Write 2-3 sentences interpreting what the coefficients tell us about each feature's impact on house prices.

2. Evaluate Model Quality

- a) Make predictions on the test set
 - b) Calculate and display the following metrics:
 - Mean Squared Error (MSE)
 - Root Mean Squared Error (RMSE)
 - R^2 Score (coefficient of determination)
 - c) Analysis: Write 2-3 sentences interpreting these evaluation metrics and assessing the model's performance.
3. Predict for the Excluded House
- a) Find the house with id '1925069082' in the original dataset
 - b) Extract the feature values for this house
 - c) Use your trained model to predict the price of this house
 - d) Display the actual price of this house
 - e) Calculate the prediction error:
 - Absolute error
 - Percentage error
 - f) Analysis: Write 2-3 sentences evaluating how accurate your model's prediction was for this specific house.

- End of Tutorial 4 -