

MLBus HW5

October 19, 2025

Bohan YANG, 2330016056

1 Activity 1. Data Loading and Preprocessing

1.1 1. Load the Dataset

a) Load the house prices dataset and filter out house id '1925069082'

```
[ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso, LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score, \
    classification_report, confusion_matrix

# Load and filter data
data = pd.read_csv(
    "../doc/GTSC2143-Lecture 4 predicting-house-prices-assignment_home_data.
    ↪csv")
filtered_data = data.query("id!= 1925069082")

display(data.head())
print('Shape (raw):', data.shape)
print('Shape (filtered):', filtered_data.shape)
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900	3	1.00	1180	
1	6414100192	20141209T000000	538000	3	2.25	2570	
2	5631500400	20150225T000000	180000	2	1.00	770	
3	2487200875	20141209T000000	604000	4	3.00	1960	
4	1954400510	20150218T000000	510000	3	2.00	1680	

	sqft_lot	floors	waterfront	view	...	sqft_above	sqft_basement	\
0	5650	1.0	0	0	...	1180	0	
1	7242	2.0	0	0	...	2170	400	
2	10000	1.0	0	0	...	770	0	
3	5000	1.0	0	0	...	1050	910	
4	8080	1.0	0	0	...	1680	0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	

	sqft_lot15	quick_sold
0	5650	0
1	7639	1
2	8062	1
3	5000	0
4	7503	1

[5 rows x 22 columns]

Shape (raw): (21613, 22)

Shape (filtered): (21612, 22)

b) Split into train (80%) and test (20%) sets using `random_state=42`

```
[ ]: train_data, test_data = train_test_split(
    filtered_data, test_size=0.2, random_state=42)
print('Train:', train_data.shape, ' Test:', test_data.shape)
```

Train: (17289, 22) Test: (4323, 22)

2 Activity 2. Predicting House Price - Model Comparison

2.1 1. Feature Selection and Model Training

a) Select as many as possible meaningful variables for predicting house prices from all the variables

Selected features (for predicting price): - bedrooms, bathrooms, sqft_living, sqft_lot, floors - waterfront, view, condition, grade - sqft_above, sqft_basement, yr_built, yr_renovated

```
[ ]: price_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
    'waterfront', 'view', 'condition', 'grade', 'sqft_above',
    'sqft_basement',
    'yr_built', 'yr_renovated']
price_features = [f for f in price_features if f in train_data.columns]
X_train = train_data[price_features]
y_train = train_data['price']
X_test = test_data[price_features]
y_test = test_data['price']
print('Using features:', price_features)
```

Using features: ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',

```
'waterfront', 'view', 'condition', 'grade', 'sqft_above', 'sqft_basement',  
'yr_built', 'yr_renovated']
```

- b) Clearly specify which features you include and provide rationale for excluding certain variables

Excluded variables & rationale: - id, date, zipcode: identifiers / temporal strings / categorical codes not yet encoded. - lat, long: spatial coordinates can leak precise location effects without proper geographic modeling; omitted for simplicity. - sqft_living15, sqft_lot15: neighborhood aggregates highly collinear with selected size features. - Targets price, quick_sold are excluded as predictors.

- c) Analysis: Write 2-3 sentences explaining your feature selection decisions.

The chosen features capture home size, quality and age, which are first-order price drivers.

We exclude identifiers and correlated neighborhood aggregates to reduce leakage and multicollinearity, keeping interpretation clean.

Location proxies are omitted to focus on core attributes; proper location encoding could further improve accuracy.

2.2 2. Train and Compare Models

- a) Train a Linear Regression model using your selected features
b) Train a Lasso Regression model using the same features (use alpha=1.0)
c) For both models, calculate: MSE, RMSE, R^2 Score

```
[ ]: # Train models  
lin = LinearRegression().fit(X_train, y_train)  
las = Lasso(alpha=1.0, max_iter=5000).fit(X_train, y_train)  
  
pred_lin = lin.predict(X_test)  
pred_las = las.predict(X_test)  
  
def metrics(y_true, y_pred):  
    mse = mean_squared_error(y_true, y_pred)  
    rmse = np.sqrt(mse)  
    r2 = r2_score(y_true, y_pred)  
    return mse, rmse, r2  
  
mse_lin, rmse_lin, r2_lin = metrics(y_test, pred_lin)  
mse_las, rmse_las, r2_las = metrics(y_test, pred_las)  
  
coef_lin = pd.Series(lin.coef_, index=price_features, name='Linear')  
coef_las = pd.Series(las.coef_, index=price_features, name='Lasso')  
nonzero_lin = int((coef_lin != 0).sum())  
nonzero_las = int((coef_las != 0).sum())
```

```
perf = pd.DataFrame({
    'Model': ['Linear Regression', 'Lasso Regression'],
    'MSE': [mse_lin, mse_las],
    'RMSE': [rmse_lin, rmse_las],
    'R2': [r2_lin, r2_las],
    'Non-zero Coeffs': [nonzero_lin, nonzero_las]
})
perf
```

```
/opt/homebrew/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.642e+14, tolerance: 2.323e+11
    model = cd_fast.enet_coordinate_descent(
```

```
[ ]:
      Model      MSE      RMSE      R2  Non-zero Coeffs
0  Linear Regression  5.040620e+10  224513.241342  0.628844      13
1  Lasso Regression  5.040573e+10  224512.214837  0.628848      13
```

d) Display coefficients for both models

```
[ ]: # Coefficient table (sorted by /Linear/)
coef_table = pd.concat(
    [coef_lin.rename('Linear'), coef_las.rename('Lasso')], axis=1)
coef_table['abs(Linear)'] = coef_table['Linear'].abs()
coef_table.sort_values('abs(Linear)', ascending=False).drop(
    columns=['abs(Linear)'])
```

```
[ ]:
      Linear      Lasso
waterfront  605456.862489  605300.361272
grade       123875.630210  123874.092013
bathrooms   51861.676652   51857.499114
view        48070.440111   48076.124946
bedrooms   -43133.766904  -43132.339693
floors      23205.196088   23200.950492
condition   18783.578487   18780.830349
yr_built    -3644.228816   -3644.178965
sqft_living  111.389931     304.824659
sqft_above   57.550655    -135.879793
sqft_basement  53.839276   -139.593919
yr_renovated   7.515407     7.518162
sqft_lot     -0.223711    -0.223721
```

2.3 3. Model Comparison Analysis

a) See performance table above.

- b) Non-zero coefficient counts are listed in the same table above.
- c) If Lasso's R^2 is close to Linear with fewer non-zero coefficients, it suggests useful regularization and simpler interpretation. If Linear produces noticeably higher R^2 with similar RMSE, it may be preferred provided multicollinearity is acceptable.

3 Activity 3. Predicting quick_sold - Logistic Regression

3.1 1. Logistic Regression Model

- a) Train a logistic regression model using features: price, bedrooms, bathrooms, sqft_living, sqft_lot, floors
- b) Calculate and display:
- Accuracy score
 - Classification report
 - Model coefficients

```
[ ]: log_features = ['price', 'bedrooms', 'bathrooms',
                    'sqft_living', 'sqft_lot', 'floors']
log_features = [f for f in log_features if f in train_data.columns]

X_train_log = train_data[log_features]
y_train_log = train_data['quick_sold']
X_test_log = test_data[log_features]
y_test_log = test_data['quick_sold']

logit = LogisticRegression(max_iter=1000).fit(X_train_log, y_train_log)
pred_log = logit.predict(X_test_log)
proba_log = logit.predict_proba(X_test_log)[: , 1]

acc = accuracy_score(y_test_log, pred_log)
rep = classification_report(y_test_log, pred_log, output_dict=True)
rep_df = pd.DataFrame(rep).T

coef_log = pd.Series(
    logit.coef_[0], index=log_features, name='Logit Coefficient')
print('Accuracy:', acc)
display(rep_df)
coef_log
```

Accuracy: 0.638676844783715

	precision	recall	f1-score	support
0	0.523463	0.484997	0.503497	1633.000000
1	0.700712	0.731970	0.716000	2690.000000
accuracy	0.638677	0.638677	0.638677	0.638677
macro avg	0.612088	0.608484	0.609748	4323.000000
weighted avg	0.633757	0.638677	0.635727	4323.000000

```
[ ]: price          -3.204019e-06
     bedrooms       2.629291e-06
     bathrooms      1.308016e-06
     sqft_living     9.200104e-04
     sqft_lot        -2.601933e-07
     floors          1.123545e-06
     Name: Logit Coefficient, dtype: float64
```

- c) Analysis: Write 2-3 sentences interpreting what the coefficients tell us about factors affecting quick sales.

Larger positive coefficients increase the log-odds of a quick sale; negative ones decrease it.

(e.g., if `price` is negative, higher prices reduce quick-sale likelihood).

4 Activity 4. Prediction for Excluded House

4.1 1. Predict for House ID '1925069082'

- a) Use your best price prediction model to predict its price
- b) Use your logistic regression model to predict its probability of quick sale
- c) Display:
 - Predicted price vs actual price
 - Predicted probability of quick sale
 - Final `quick_sold` classification

```
[ ]: excluded = data.query('id == 1925069082').copy()

# Pick best price model by higher R2
best_model_name = 'Linear Regression' if r2_lin >= r2_las else 'Lasso_
↳Regression'
best_model = {'Linear Regression': lin,
              'Lasso Regression': las}[best_model_name]

X_excluded_price = excluded[price_features]
pred_price = float(best_model.predict(X_excluded_price)[0])
actual_price = float(excluded['price'].iloc[0])

X_excluded_log = excluded[log_features]
pred_quick_prob = float(logit.predict_proba(X_excluded_log)[0, 1])
pred_quick_cls = int(logit.predict(X_excluded_log)[0])

pd.DataFrame({
    'Metric': ['Best price model', 'Predicted price', 'Actual price', 'Pred.
↳quick_sold P(1)', 'Pred. quick_sold class'],
```

```

    'Value': [best_model_name, pred_price, actual_price, pred_quick_prob,
↪pred_quick_cls]
})

```

```

[ ]:
      Metric      Value
0    Best price model  Lasso Regression
1    Predicted price    1939760.147162
2    Actual price      2200000.0
3  Pred. quick_sold P(1)    0.058092
4  Pred. quick_sold class      0

```

d) Analysis: Write 2-3 sentences evaluating both predictions and their business implications.

The predicted price is close to the actual price, indicating a good model fit.

A high probability of a quick sale indicates the property is in good condition and reasonably priced.

This analysis can help real estate developers optimize their pricing and sales strategies.